




CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

ETH zürich



Automatic Building & Continuous Validation CSCS Roadmap

Guilherme Peretti-Pezzi
User Engagement and Support

11.06.2015 hpc-ch forum

Outline



- Background
- Automatic Building tools for HPC
- Continuous Validation
- CSCS preliminary use cases
- Hands on Demo

Evolution of requirements for scientific applications

- A typical requirement workflow in a galaxy not so far away
 - Application: Fortran / C
 - Numerical library: Fortran
 - Parallel API: MPI
- Nowadays
 - Application
 - Fortran / C / C++ / Python
 - Numerical libraries
 - Fortran / C
 - Python + scipy, numpy, matplotlib, h5py, mpi4py, fipy, netCDF4 + (...)
 - R + Modules
 - I/O libraries HDF, netcdf
 - Parallel API
 - MPI / PGAS / CUDA / OpenMP / OpenACC

Current goals @ CSCS

- Minimize the maintenance overhead
 - Number of supported applications is constantly increasing
 - Upgrades are becoming a headache
- Provide a uniform set of libraries and applications
 - On ALL of our systems
- Improve QoS by adopting continuous validation techniques
 - Unit/Non-regression testing
- Provide a continuous integration platform to the users
 - Ease the development and improve quality of scientific software
- Challenges
 - Scale (number of systems, users and applications)
 - Implement consistent procedures within and across large teams

Standard package management tools

- Popular examples
 - Yum (RPM)
 - Deb
 - Homebrew
 - Portage
 - pkgsrc
 - Nix
- Known issues
 - Lack of isolation
 - Not adapted for handling multiple versions of the same software
 - Limited support for automatic building
 - Reproducibility
 - Lack of integration with modules environment
 - Error-prone (manual/interactive steps)

Smithy

- ORNL
- Ruby-based
- 100+ of supported packages
- Pros
 - Simplicity
 - One file describes all the building process
 - Architecture agnostic
- Cons
 - Builds are described as a program-like syntax
 - Lack of community / users outside ORNL

Spack

- LLNL
- Python based
- 100+ of supported packages
- Pros
 - DSL for describing dependencies/workflow
 - Dependencies can be described without hardcoding version number
 - Architecture agnostic
- Cons
 - Lack of community / users outside LLNL

EasyBuild

- HPC-Ugent
- Python based
- 600+ of supported packages
- Pros
 - Two-level dependency description
 - Easyconfig files for standard builds
 - Easyblocks (python code) for more complex
 - **Growing community**
 - sciCORE/UniBas, Flemish Supercomputer Centre, Julich Supercomputer Centre, Stanford Univ., Univ. of Auckland, Bayer AG, Texas A&M, IMB (Austria), Univ. of Luxembourg, Cyprus Institute
- Cons
 - One easyconfig file maps to one specific software and toolchain version
 - Thousands of files..
 - Focused on building toolchains from scratch
 - Non trivial procedure for creating/extending toolchains
 - Version 2.1.0 (03/2015) provides limited Cray support

Overview of Automatic Building tools for HPC

	Smithy	Spack	EasyBuild
Language	Ruby	Python	Python
# of packages	150+	100+	600+
High level Input file	No	Yes	Yes
Module file Generation	Yes	Yes	Yes

Testing EasyBuild 2.1.1 @ CSCS (as of 11.06.2015)

- Linux x64 clusters (pilatus, castor)
 - OK using foss-2015a toolchain and a couple of stock easyconfig files
- Cray XC30/XC40 (Piz Daint, Santis, Dora)
 - OK for stock Python-2.7.9 using CrayGNU toolchain
 - Created easyconfig file for Python-3.4.3
 - Modules scipy, numpy, matplotlib, mpi4py, virtualenv, pip, ...
- Cray MeteoSwiss (Early access)
 - Requirements: GCC 4.8.2 + MVAPICH2
 - PrgEnv-gnu not available
 - Falling back to EB gmvapich2 toolchain
 - Updated to MVAPICH 2.1
 - Next steps: Intel + PGI

Continuous Integration/Validation Tools

- Cdash
 - From Cmake developers
- Bamboo
 - From Atlassian
- Jenkins
 - Formerly Hudson
 - Widely used

1st use case @ CSCS: Automatic builds

- Idea:
 - Regular builds (test) and automatic install on all systems
- EasyBuild
 - Cray: Python/2.7.9 + 3.4.3 + modules,
 - Gnu: Bison/3.0.2, flex/2.5.39, Cmake/3.0.0, HDF5/1.8.15, netCDF/4.3.3.1
- Jenkins

<div>All EasyBuild +</div>					
S	W	Name ↓	Last Success	Last Failure	Last Duration
		CrayGNU Unit Test on Daint	3 days 0 hr - #1	N/A	13 min 
		EasyBuild Apps on daint	2 days 23 hr - #9	3 days 0 hr - #7	49 min 
		EasyBuild apps on Pilatus	2 days 22 hr - #7	N/A	1 hr 22 min 
		EasyBuild Apps on Santis	3 days 0 hr - #4	N/A	2 min 19 sec 

2nd use case @ CSCS: DCA++

- Idea
 - Provide required libraries with EasyBuild
 - Regularly run builds, unit and performance tests
- EasyBuild (TBD)
 - Libraries: HDF5, NFFT, SPGLIB
- Google Mock + Ctest
 - Unit testing (at application level)
 - Sequential
 - Parallel (MPI)
- Jenkins
 - Github repository integration
 - Automatic build of App + dependencies
 - Automatic execution of unit test
 - Triggered by git commits

3rd Use case @ CSCS: Non-regression suite

- Current state
 - Bash scripts for simple tests
 - Hardware/System
 - Compilation environment sanity
 - Performance
 - Manually triggered
- Next steps
 - EasyBuild + Jenkins integration
 - Extend suite for testing for more complex applications
 - Automatic building
 - To ease the work when upgrading or performing maintenance
 - Monitor “user experience”
 - Periodic performance tests of (key) applications
 - For example: FS and Slurm response times

EB + Jenkins Workflow @ Piz Daint


- Setup EasyBuild
 - `$ source $APPS/easybuild/setup.sh`
- Create a new EasyConfig file using CrayGNU toolchain
- Build and test application
 - `$ eb newApp.eb -r`
- Move .eb file to shared folder for nightly builds
 - `/apps/common/easybuild/cscs_easyconfigs/`
 - CRAY_XC30 (for Piz Daint, Piz Dora and Santis)
 - GNU for non-Cray (Pilatus, Castor, ...)
- Day 2
 - Check Jenkins build output on all systems
 - If OK copy new .eb file to automatic deployment folder



CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

ETH zürich



Thank you for your attention.