

# EasyBuild 2.3 on Cray Linux Environment

(with some CSCS specifics)

Petar Forai

# Requirements

- Tcl/C Environment Modules **3.2.10** or  $\geq$  Lmod 5.6.3
- Lmod can consume Cray module files!  
(Thanks and Kudos to Robert McLay for this!)
- CLE 5.x with  $\geq$ **craype-2.0** because we only have XC40 systems for developing EB

# Current Status as of 2.3

- Some big apps work and most of the machinery is there
- Marked as 'experimental', ie current interfaces and functionality could be changed in future (waiting for feedback from you guys!)
- Needs —**experimental** or environment variable **\$EASYBUILD\_EXPERIMENTAL** defined to get EB to make experimental parts of it available
- Updated status page at

<https://github.com/hpcugent/easybuild/wiki/EasyBuild-on-Cray>

# EasyBuild approach for toolchains

- Toolchain definitions are composition of ‘full stack’ development environments (compilers, MPI, BLAS/LAPACK, FFTW, etc.)
- Toolchain code in framework exposes things like `$CC`, `$CXX`, `$FTN`, `$CFLAGS`, `$CXXFLAGS`, etc. when building
- There’s also API in the toolchains to construct compiler call strings - mostly used in complex blocks
- *EB paradigm is to have reproducible software builds and EB archives that by pinning all versions of all of the TC’s dependencies*

# Mapping Cray PrgEnv(s) to EasyBuild

- EB offers toolchains for `PrgEnv-{gnu|intel|cray}`, they just load `PrgEnv-{gnu|intel|cray}` in a specific version (same as toolchain), and `fftw` (*system defaults!*)- which makes this *similar* to other TCs
- Named `Cray{Gnu|Intel|CCE}-<PrgEnv-version>` Ie `CrayGNU-5.2.40` uses `PrgEnv-gnu/5.2.40`
- Install the resulting configs to have the toolchains available for building.
- If a new version of `PrgEnv-gnu` is provided on the system, create a new EB config for this and install as you would any other piece of software.

```
easyblock = 'Toolchain'

name = 'CrayGNU'
version = '5.2.40'

homepage = '(none)'
description = """Toolchain using Cray compiler wrapper, using PrgEnv-gnu module."""

toolchain = {'name': 'dummy', 'version': 'dummy'}

dependencies = [
    ('PrgEnv-gnu/' + version, EXTERNAL_MODULE), # also loads cray-libsci, dmi, atp, etc.
    ('fftw/3.3.4.2', EXTERNAL_MODULE),
]

moduleclass = 'toolchain'
```

# Working with EB on Cray

## Environment setup

- On non-Crays EB builds everything from scratch and provide the appropriate modules
- Please note that we rely on the environment to provide the configuration for MPI. Consider this if you run `module purge && eb foo-CrayGNU-5.2.40.eb`!
- Setup scripts for CSCS just runs `module swap PrgEnv-cray PrgEnv-gnu`
- This way the environment is bootstrapped through system defaults, only compiler swapped. MPI and anything else is there already and system module default versions are loaded

# Working with EB on Cray

## External Modules

- External modules supported from 2.1 via meta data description
- Maps software available on system to software available within EasyBuild, useful to integrate Cray provided software in EB (HDF5, third party libraries, etc.)
- Example <https://gist.github.com/pforai/d0f25022b7925792f9f7>
- CSCS's global EB setup for metadata is `/apps/common/easybuild/cray_external_modules_metadata.cfg`

# Working with EB on Cray

## External Modules Meta Data

- Normally EasyBuild builds everything and generates modules with known and defined structure
- If piece of software exists in EasyBuild then supplying external data via a central config file is the currently supported approach
- Map module name and version into EB software name space
- Example for Cray HDF5 parallel and serial into EB HDF5 build and NetCDF

```
[cray-netcdf/4.3.2]
name = netCDF,netCDF-Fortran
version = 4.3.2, 4.3.2
prefix = NETCDF_DIR
```

```
[cray-hdf5/1.8.13]
name = HDF5
version = 1.8.13
prefix = HDF5_DIR
```

```
[cray-hdf5-parallel/1.8.13]
name = HDF5
version = 1.8.13
prefix = HDF5_DIR
```

```
dependencies = [
    ('cray-netcdf/4.3.2', EXTERNAL_MODULE),
    ('cray-hdf5-parallel/1.8.13', EXTERNAL_MODULE),
]
```



# Working with EB on Cray CPU targeting

- Cray's wrappers accept different backend code generators options/targets through environment variables
- “Cross compilation” works by loading craype-haswell or craype-ivybridge modules or others if machine has different CPUs in login nodes vs backend nodes via the **—optarch=<target>** flag or by setting **EASYBUILD\_OPTARCH=<craype-target>**
- This must be specified otherwise EB will refuse to start

# Working with EB on Cray Linking

- EasyBuild controls linking by exposing `$CRAY_LINK_TYPE` into the environment while building, otherwise that variable is undefined and default wrapper behaviour applies
- default is to do static linking
- Controlled through toolchain options in EasyBuild configs

```
easyblock = 'ConfigureMake'

name = 'zlib'
version = '1.2.8'

toolchain = {'name': 'CrayGNU', 'version': '5.2.40'}
toolchainopts = {'pic': True, 'dynamic': True}
```

# Working with EB on Cray

## Other toolchain options

- Get verbose builds by setting verbose TC option

```
...  
toolchain = {'name': 'CrayGNU', 'version': '5.2.40'}  
toolchainopts = {'pic': True, 'dynamic': True, 'verbose': True}  
...
```

- Enable fine grained multi threading support in MPICH

```
toolchainopts = {'pic': True, 'dynamic': True, 'mpich-mt': True,}
```

# Implementation Effort up until 2.3

- Less than 400 SLOC (including external modules support, generic cray support, GNU/Intel/CCE specifics)
- in `easybuild/toolchains/craype.py` (base class)
- in `easybuild/framework/easyconfig/easyconfig.py`
- in `easybuild/tools/toolchain.py`

# Potential Caveats with EB

- EB assumes it can execute tests where it is running (can be disabled)
- Executing and/or benchmarking build artefacts (ie numpy for np.dot() needs to run <1s) after the build for sanity checking
- Linking is controlled by setting `$CRAYPE_LINK_TYPE` (no call out to `-shared` or `-dynamic`)

# What works

- Some large apps work as expected, like CP2K, GROMACS, WRF, Python with numpy and scipy linked against libsci and cray-mpich for mpi4py (using cray supplied software as dependencies)
- All login nodes software (git, mercurial, Autotools, cURL, Qt4/5, ...) that just relies on dummy compiler or
- Simple to medium simple software for BioInformatics, Math libraries, etc.

# What is missing

- Most work has gone into getting PrgEnv-gnu, obviously the other PrgEnv- modules will need some more work, mostly in blocks and of course configs
- Boost and some other major dependencies that needs custom blocks
- Generating pkg-config .pc files at end of build to have the wrappers set proper -I and -L -I for 3rd party software (WIP) (workaround available)
- Cleaning ups configs and blocks - see <https://github.com/hpcugent/easybuild-easyconfigs/search?utf8=%E2%9C%93&q=CFLAGS> (there are assumptions in (some) places that the compilers running accept GCC flags)

# Open Questions

- What to do with libsci? EB normally pins all the version of all dependencies. PrgEnv-foo just loads system default for libsci (unlike other some deps in PrgEnv modules)
- Same for GCC!
- On one machine this means that CrayGNU-5.2.25 contains ie libsci 13.0.3 and on another it is 13.0.1
- Current assumption is 'site knows what it's doing and is checking for compatibility'
- Should it be pin and if so how to parametrise?
- Provide feedback for EB built software in terms of performance vs your hand compiled stacks (like for large CP2K runs, GROMACs, etc.)?



# Open Questions II

- how to correctly deal with MPI: which modules should be loaded? EB Cray toolchain should load an MPI module? Current assumption is 'site knows what it's doing and is checking for compatibility'
- Do you "unwrap" the drivers? Do you "unwrap" the drivers? I.e. load PrgEnv-gnu/x.y.z but set \$CC to gcc?
- Want to approach Cray Inc. at Super Computing 2015 and present this to them. CSC (Olli Pekka Lehto) was approaching them before about EasyBuild.